

## Query Optimization Example

- Sailors (sid, sname, rating, age)
- Boats(bid, bname, color)
- Reserves(sid, bid, day, rname)
- Query:  

```
SELECT S.sid, S.sname, S.age
FROM   Sailors S, Boats B, Reserves R
WHERE  B.bid = R.rid AND B.bid = R.bid AND
       B.color = "Red" AND S.age < 30;
```
- Reserves has 1000 pages, 10 tuples/page
- Sailors has 500 pages, 50 tuples/page
- Boats has 160 pages, 10 tuples/page
- Data is evenly distributed (assumption)

CPSC 421, 2009

1

## Steps

- Query optimization steps:
  - Translate SQL Query to Relational Algebra Query
  - Create a query tree
  - Create left-deep alternative trees
  - Create query plans for our trees
  - Estimate costs of plans
  - Pick best one

CPSC 421, 2009

2

## Step 1

- Translate SQL Query to Relational Algebra Query


$$\pi_{\text{sid, sname, age}}(\sigma_{\text{age} < 30}(\sigma_{\text{color} = \text{"Red"}}(\sigma_{\text{bid} = \text{bid}}(\mathbf{B} \times \sigma_{\text{sid} = \text{sid}}(\mathbf{S} \times \mathbf{R}))))))$$

CPSC 421, 2009

3

## Step 2

- Create a query tree



Note our query only involves joins  
(as opposed to plain old cross-  
products) ...

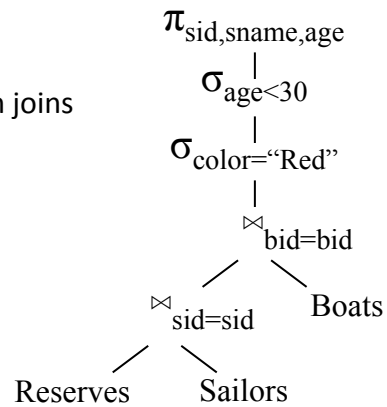
Lets draw the trees with joins

CPSC 421, 2009

4

### Step 3

- Create left-deep alternative trees
  - We replace cross products if they are really joins
  - How?
    - Using RA equivalences!
  - This is a left-deep plan with joins





CPSC 421, 2009

5

### Step 3 – More Alternatives

- Would we expect these to have different costs over previous alternative?



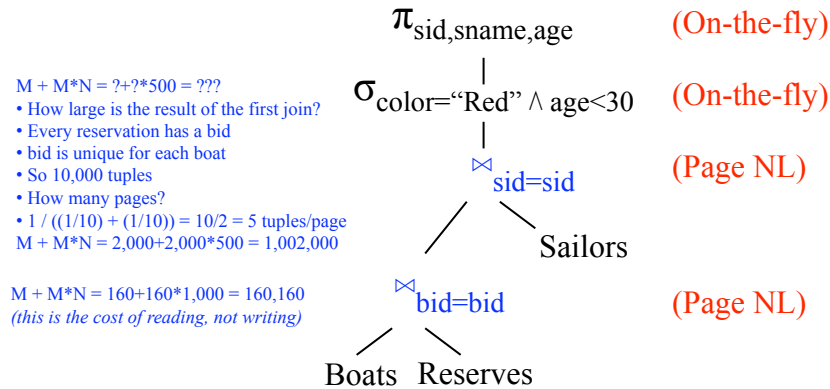
CPSC 421, 2009

6

## Step 4 & 5 – Plans and Estimates

- Lets create some plans and estimates

Total cost:  $160,160 + 1,002,000 = 1,162,160$  I/Os !!!



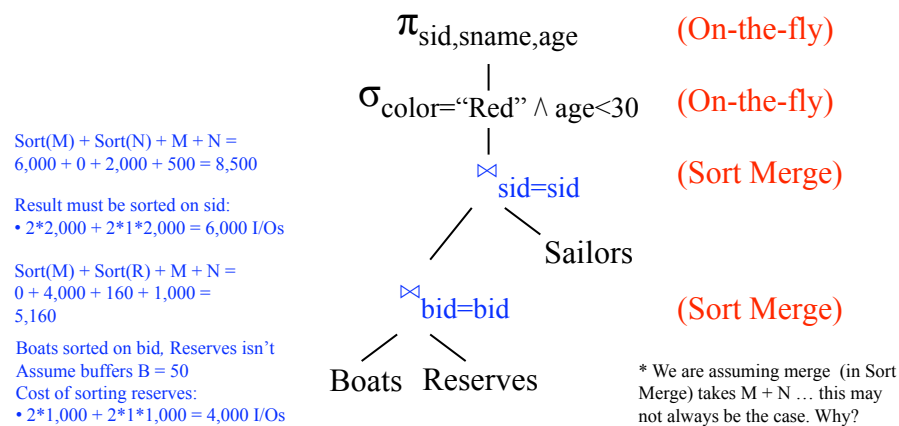
CPSC 421, 2009

7

## Step 4 & 5 – Plans and Estimates

- Assume Boats has a clustered B+ Tree on bid (key)
- Assume S has a clustered B+ Tree on sid (key)

Total cost:  $5,160 + 8,500 = 13,660$  I/Os !!!



CPSC 421, 2009

8

## Note on Sort Merge

*This is a much better plan than our previous page nested loops one*

- But it is an overestimate for Sort-Merge join!
  - In Pass 0 we read all pages then write all pages ( $2 \cdot M$ )
  - We do not need to read all pages if they are pipelined to the next join ( $1 \cdot M$ )
  - In last merge pass, we don't have to write the last set of pages since we always pipeline them to the next operator
  - We also don't have to read in the file again:
    - $\text{Sort}(M) + \text{Sort}(N)$  if  $N$  needs to be sorted
    - $\text{Sort}(M) + N$  if  $N$  is sorted

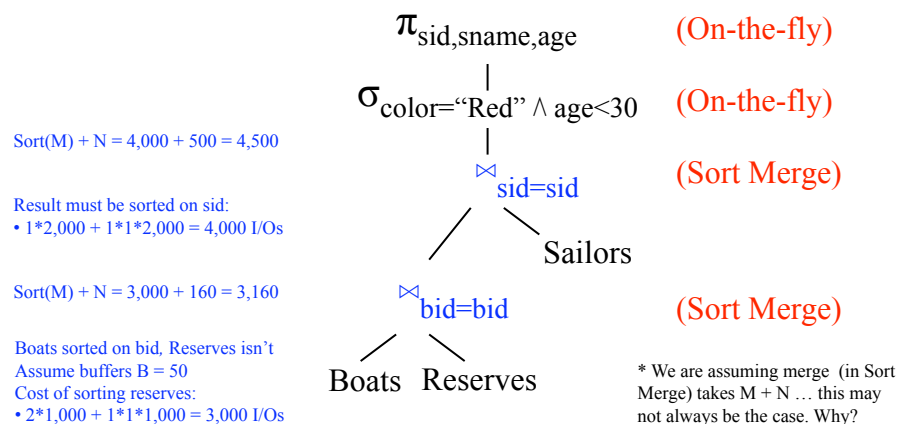
CPSC 421, 2009

9

## Step 4 & 5 – Plans and Estimates

- Assume Boats has a clustered B+ Tree on bid (key)
- Assume S has a clustered B+ Tree on sid (key)

Total cost:  $3,160 + 4,500 = 7,660$  I/Os !!!



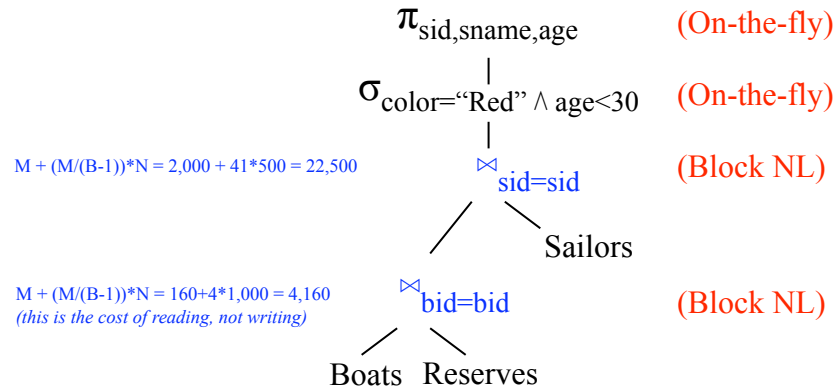
CPSC 421, 2009

10

## Step 4 & 5 – Plans and Estimates

- Lets try block nested loop join (B=50)

Total cost:  $4,160 + 22,500 = 26,660$  I/Os !!!



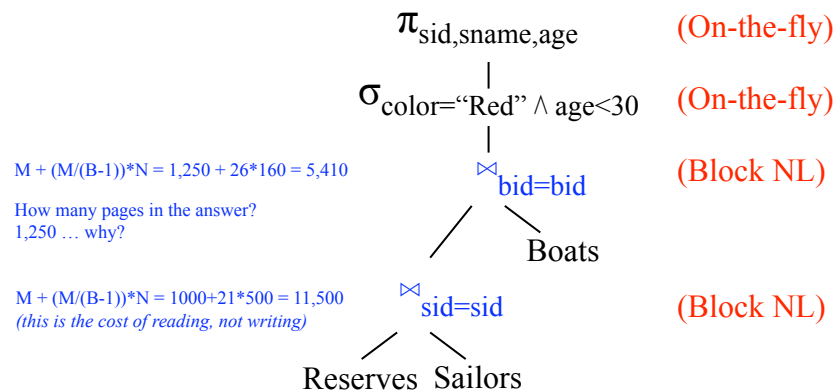
CPSC 421, 2009

11

## Step 4 & 5 – Plans and Estimates

- Can we do better?

Total cost:  $5,410 + 11,500 = 16,910$  I/Os !!!



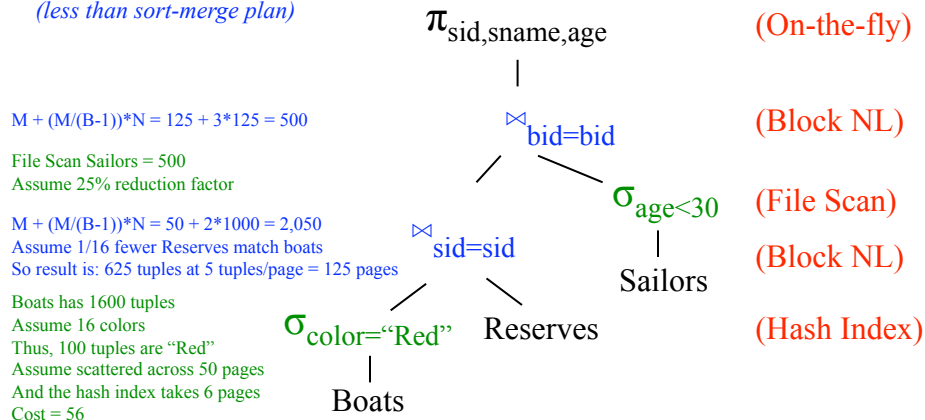
CPSC 421, 2009

12

## Step 4 & 5 – Plans and Estimates

- Can we do even better?
  - Assume Boats has a hash index on color

Total cost:  $56 + 2,050 + 500 + 500 = 3,106$  I/Os !!!  
*(less than sort-merge plan)*



CPSC 421, 2009

13

## Step 4 & 5 – Plans and Estimates

- Can we still do better?
  - Yes!
  - E.g., we could project on bid after first join ...
- Notice that by adding an index on color to Boats ...
  - We reduced our best query time in half!
  - This is one reason we talk about query optimization
  - It drives physical database design
  - Improving performance (by adding indexes, e.g.) should be driving on how query optimization works!

CPSC 421, 2009

14